In [24]:
```
import pandas as pd
import numpy as np
```

In [3]:
```
people={
    "first":["Corey","Jane","John"],
    "last":["Schafer","Doe","Doe"],
    "email":["CoreyMSchafer@gmail.com","JaneDoe@gmail.com","JohnDoe@gmail.com"]
}
```

In [4]:
```
df=pd.DataFrame(people)
```

In [5]:
```
df
```

Out[5]:

	first	last	email
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@gmail.com
2	John	Doe	JohnDoe@gmail.com

In [6]:
```
df.columns
```

Out[6]:
```
Index(['first', 'last', 'email'], dtype='object')
```

In [9]:
```
df.columns=['first_name','last_name','email']
```

In [10]:
```
df
```

Out[10]:

	first_name	last_name	email
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@gmail.com
2	John	Doe	JohnDoe@gmail.com

In [11]:
```
df.columns=[x.upper() for x in df.columns]
```

In [12]:
```
df
```

Out[12]:

	FIRST_NAME	LAST_NAME	EMAIL
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@gmail.com
2	John	Doe	JohnDoe@gmail.com

In [13]:
```
df.columns=df.columns.str.replace('_',' ')
```

In [14]:
```
df
```

Out[14]:

	FIRST NAME	LAST NAME	EMAIL
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@gmail.com
2	John	Doe	JohnDoe@gmail.com

In [15]:
```
df.columns=[x.lower() for x in df.columns]
```

In [16]:
```
df
```

Out[16]:

	first name	last name	email
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@gmail.com
2	John	Doe	JohnDoe@gmail.com

In [17]:
```
df.columns=df.columns.str.replace(' ','_')
```

In [18]:
```
df
```

Out[18]:

	first_name	last_name	email
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@gmail.com
2	John	Doe	JohnDoe@gmail.com

In [20]:
```
# want to change specific columns
df.rename(columns={'first_name':'first','last_name':'last'}, inplace=True)
```

In [21]:
```
df
```

Out[21]:

	first	last	email
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@gmail.com
2	John	Doe	JohnDoe@gmail.com

In [22]:
```
df.loc[2]=['John','Smith','JohnDoe@email.com']
```

In [23]:
```
df
```

Out[23]:

	first	last	email
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@gmail.com
2	John	Smith	JohnDoe@email.com

In [24]:
```
df.loc[2,['last','email']]=['Doe','JohnDoe@gmail.com']
```

In [25]:
```
df
```

Out[25]:

	first	last	email
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@gmail.com
2	John	Doe	JohnDoe@gmail.com

In [26]:
```
df.loc[2,'last']='Smith'
```

In [27]:
```
df
```

Out[27]:

	first	last	email
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@gmail.com
2	John	Smith	JohnDoe@gmail.com

In [28]:
```
filt=(df['email']=='JohnDoe@gmail.com')
df[filt]
```

Out[28]:

	first	last	email
2	John	Smith	JohnDoe@gmail.com

In [29]:
```
df[filt]['last']='Smith'
```

```
<ipython-input-29-5c4ea8a4e6cd>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
  df[filt]['last']='Smith'
```

In [32]:
```
df['email']=df['email'].str.lower()
```

In [33]:
```
df
```

Out[33]:

	first	last	email
0	Corey	Schafer	coreymschafer@gmail.com
1	Jane	Doe	janedoe@gmail.com
2	John	Smith	johndoe@gmail.com

In [36]:
```
#DateTime index and analysis
import pandas as pd
d_parser = lambda x: pd.datetime.strptime(x, '%Y-%m-%d %I-%p')
df = pd.read_csv(r'C:\Users\e16379\Desktop\ETH_1h.csv',parse_dates=['Date'], date
df.head()
```

```
<ipython-input-36-3e4c694d82f9>:2: FutureWarning: The pandas.datetime class is deprecated and will be removed from pandas in a future version. Import from datetime module instead.
  d_parser = lambda x: pd.datetime.strptime(x, '%Y-%m-%d %I-%p')
```

Out[36]:

	Date	Symbol	Open	High	Low	Close	Volume
0	2020-03-13 20:00:00	ETHUSD	129.94	131.82	126.87	128.71	1940673.93
1	2020-03-13 19:00:00	ETHUSD	119.51	132.02	117.10	129.94	7579741.09
2	2020-03-13 18:00:00	ETHUSD	124.47	124.85	115.50	119.51	4898735.81
3	2020-03-13 17:00:00	ETHUSD	124.08	127.42	121.63	124.47	2753450.92
4	2020-03-13 16:00:00	ETHUSD	124.85	129.51	120.17	124.08	4461424.71

In [37]:
```
df['Date'].dt.day_name()
```

Out[37]:
```
0          Friday
1          Friday
2          Friday
3          Friday
4          Friday
           ...
23669    Saturday
23670    Saturday
23671    Saturday
23672    Saturday
23673    Saturday
Name: Date, Length: 23674, dtype: object
```

In [10]:
```
df.loc[0, 'Date'].day_name()
```

Out[10]:
```
'Friday'
```

In [45]:
```
 df['Date'].dt.day_name()
```

Out[45]:
```
0          Friday
1          Friday
2          Friday
3          Friday
4          Friday
           ...
23669    Saturday
23670    Saturday
23671    Saturday
23672    Saturday
23673    Saturday
Name: Date, Length: 23674, dtype: object
```

In [11]:
```
df['DayOfWeek']=df['Date'].dt.day_name()
```

In [12]:
```
df
```

Out[12]:

	Date	Symbol	Open	High	Low	Close	Volume	DayOfWeek
0	2020-03-13 20:00:00	ETHUSD	129.94	131.82	126.87	128.71	1940673.93	Friday
1	2020-03-13 19:00:00	ETHUSD	119.51	132.02	117.10	129.94	7579741.09	Friday
2	2020-03-13 18:00:00	ETHUSD	124.47	124.85	115.50	119.51	4898735.81	Friday
3	2020-03-13 17:00:00	ETHUSD	124.08	127.42	121.63	124.47	2753450.92	Friday
4	2020-03-13 16:00:00	ETHUSD	124.85	129.51	120.17	124.08	4461424.71	Friday
...	...	...	...	...	...	...	...	...
23669	2017-07-01 15:00:00	ETHUSD	265.74	272.74	265.00	272.57	1500282.55	Saturday
23670	2017-07-01 14:00:00	ETHUSD	268.79	269.90	265.00	265.74	1702536.85	Saturday
23671	2017-07-01 13:00:00	ETHUSD	274.83	274.93	265.00	268.79	3010787.99	Saturday
23672	2017-07-01 12:00:00	ETHUSD	275.01	275.01	271.00	274.83	824362.87	Saturday
23673	2017-07-01 11:00:00	ETHUSD	279.98	279.99	272.10	275.01	679358.87	Saturday

23674 rows × 8 columns

In [63]:
```
df['Date'].min()
```

Out[63]:
```
Timestamp('2017-07-01 11:00:00')
```

In [64]:
```
df['Date'].max()
```

Out[64]:
```
Timestamp('2020-03-13 20:00:00')
```

In [65]:
```
df['Date'].max() - df['Date'].min()
```

Out[65]:
```
Timedelta('986 days 09:00:00')
```

In [13]:
```
filt = (df['Date'] >= pd.to_datetime('2019-01-01')) & (df['Date'] < pd.to_datetim
df.loc[filt]
```

Out[13]:

	Date	Symbol	Open	High	Low	Close	Volume	DayOfWeek
1749	2019-12-31 23:00:00	ETHUSD	128.33	128.69	128.14	128.54	440678.91	Tuesday
1750	2019-12-31 22:00:00	ETHUSD	128.38	128.69	127.95	128.33	554646.02	Tuesday
1751	2019-12-31 21:00:00	ETHUSD	127.86	128.43	127.72	128.38	350155.69	Tuesday
1752	2019-12-31 20:00:00	ETHUSD	127.84	128.34	127.71	127.86	428183.38	Tuesday
1753	2019-12-31 19:00:00	ETHUSD	128.69	128.69	127.60	127.84	1169847.84	Tuesday
...	...	...	...	...	...	...	...	...
10504	2019-01-01 04:00:00	ETHUSD	130.75	133.96	130.74	131.96	2791135.37	Tuesday
10505	2019-01-01 03:00:00	ETHUSD	130.06	130.79	130.06	130.75	503732.63	Tuesday
10506	2019-01-01 02:00:00	ETHUSD	130.79	130.88	129.55	130.06	838183.43	Tuesday
10507	2019-01-01 01:00:00	ETHUSD	131.62	131.62	130.77	130.79	434917.99	Tuesday
10508	2019-01-01 00:00:00	ETHUSD	130.53	131.91	130.48	131.62	1067136.21	Tuesday

8760 rows × 8 columns

In [38]:
```
df.set_index('Date', inplace=True)
```

In [39]:
```
df
```

Out[39]:

Date	Symbol	Open	High	Low	Close	Volume
2020-03-13 20:00:00	ETHUSD	129.94	131.82	126.87	128.71	1940673.93
2020-03-13 19:00:00	ETHUSD	119.51	132.02	117.10	129.94	7579741.09
2020-03-13 18:00:00	ETHUSD	124.47	124.85	115.50	119.51	4898735.81
2020-03-13 17:00:00	ETHUSD	124.08	127.42	121.63	124.47	2753450.92
2020-03-13 16:00:00	ETHUSD	124.85	129.51	120.17	124.08	4461424.71
...	...	...	...	...	...	...
2017-07-01 15:00:00	ETHUSD	265.74	272.74	265.00	272.57	1500282.55
2017-07-01 14:00:00	ETHUSD	268.79	269.90	265.00	265.74	1702536.85
2017-07-01 13:00:00	ETHUSD	274.83	274.93	265.00	268.79	3010787.99
2017-07-01 12:00:00	ETHUSD	275.01	275.01	271.00	274.83	824362.87
2017-07-01 11:00:00	ETHUSD	279.98	279.99	272.10	275.01	679358.87

23674 rows × 6 columns

In [83]:
```
df
```

Out[83]:

Date	Symbol	Open	High	Low	Close	Volume
2020-03-13 20:00:00	ETHUSD	129.94	131.82	126.87	128.71	1940673.93
2020-03-13 19:00:00	ETHUSD	119.51	132.02	117.10	129.94	7579741.09
2020-03-13 18:00:00	ETHUSD	124.47	124.85	115.50	119.51	4898735.81
2020-03-13 17:00:00	ETHUSD	124.08	127.42	121.63	124.47	2753450.92
2020-03-13 16:00:00	ETHUSD	124.85	129.51	120.17	124.08	4461424.71
...	...	...	...	...	...	...
2017-07-01 15:00:00	ETHUSD	265.74	272.74	265.00	272.57	1500282.55
2017-07-01 14:00:00	ETHUSD	268.79	269.90	265.00	265.74	1702536.85
2017-07-01 13:00:00	ETHUSD	274.83	274.93	265.00	268.79	3010787.99
2017-07-01 12:00:00	ETHUSD	275.01	275.01	271.00	274.83	824362.87
2017-07-01 11:00:00	ETHUSD	279.98	279.99	272.10	275.01	679358.87

23674 rows × 6 columns

In [85]:
```
df.loc['2020']
```

Out[85]:

Date	Symbol	Open	High	Low	Close	Volume
2020-03-13 20:00:00	ETHUSD	129.94	131.82	126.87	128.71	1940673.93
2020-03-13 19:00:00	ETHUSD	119.51	132.02	117.10	129.94	7579741.09
2020-03-13 18:00:00	ETHUSD	124.47	124.85	115.50	119.51	4898735.81
2020-03-13 17:00:00	ETHUSD	124.08	127.42	121.63	124.47	2753450.92
2020-03-13 16:00:00	ETHUSD	124.85	129.51	120.17	124.08	4461424.71
...	...	...	...	...	...	...
2020-01-01 04:00:00	ETHUSD	129.57	130.00	129.50	129.56	702786.82
2020-01-01 03:00:00	ETHUSD	130.37	130.44	129.38	129.57	496704.23
2020-01-01 02:00:00	ETHUSD	130.14	130.50	129.91	130.37	396315.72
2020-01-01 01:00:00	ETHUSD	128.34	130.14	128.32	130.14	635419.40
2020-01-01 00:00:00	ETHUSD	128.54	128.54	128.12	128.34	245119.91

1749 rows × 6 columns

In [86]:
```
df.loc['2020-01':'2020-02']
```

Out[86]:

Date	Symbol	Open	High	Low	Close	Volume
2020-02-29 23:00:00	ETHUSD	223.35	223.58	216.83	217.31	1927939.88
2020-02-29 22:00:00	ETHUSD	223.48	223.59	222.14	223.35	535998.57
2020-02-29 21:00:00	ETHUSD	224.63	225.14	222.74	223.48	561158.03
2020-02-29 20:00:00	ETHUSD	225.31	225.33	223.50	224.63	511648.65
2020-02-29 19:00:00	ETHUSD	225.09	225.85	223.87	225.31	1250856.20
...	...	...	...	...	...	...
2020-01-01 04:00:00	ETHUSD	129.57	130.00	129.50	129.56	702786.82
2020-01-01 03:00:00	ETHUSD	130.37	130.44	129.38	129.57	496704.23
2020-01-01 02:00:00	ETHUSD	130.14	130.50	129.91	130.37	396315.72
2020-01-01 01:00:00	ETHUSD	128.34	130.14	128.32	130.14	635419.40
2020-01-01 00:00:00	ETHUSD	128.54	128.54	128.12	128.34	245119.91

1440 rows × 6 columns

In [88]:
```
df.loc['2020-01':'2020-02']['Close'].mean()
```

Out[88]:
```
195.16559027777814
```

In [90]:
```
df.loc['2020-01-01']['High'].max()
```

Out[90]:
```
132.68
```

In [16]:
```
highs = df['High'].resample('D').max()
```

In [17]:
```
df
```

Out[17]:

Date	Symbol	Open	High	Low	Close	Volume	DayOfWeek
2020-03-13 20:00:00	ETHUSD	129.94	131.82	126.87	128.71	1940673.93	Friday
2020-03-13 19:00:00	ETHUSD	119.51	132.02	117.10	129.94	7579741.09	Friday
2020-03-13 18:00:00	ETHUSD	124.47	124.85	115.50	119.51	4898735.81	Friday
2020-03-13 17:00:00	ETHUSD	124.08	127.42	121.63	124.47	2753450.92	Friday
2020-03-13 16:00:00	ETHUSD	124.85	129.51	120.17	124.08	4461424.71	Friday
...	...	...	...	...	...	...	...
2017-07-01 15:00:00	ETHUSD	265.74	272.74	265.00	272.57	1500282.55	Saturday
2017-07-01 14:00:00	ETHUSD	268.79	269.90	265.00	265.74	1702536.85	Saturday
2017-07-01 13:00:00	ETHUSD	274.83	274.93	265.00	268.79	3010787.99	Saturday
2017-07-01 12:00:00	ETHUSD	275.01	275.01	271.00	274.83	824362.87	Saturday
2017-07-01 11:00:00	ETHUSD	279.98	279.99	272.10	275.01	679358.87	Saturday

23674 rows × 7 columns

In [18]:
```
highs['2020-01-01']
```

Out[18]:
```
132.68
```

In [19]:
```
%matplotlib inline
highs.plot()
```

Out[19]:
```
<AxesSubplot:xlabel='Date'>
```

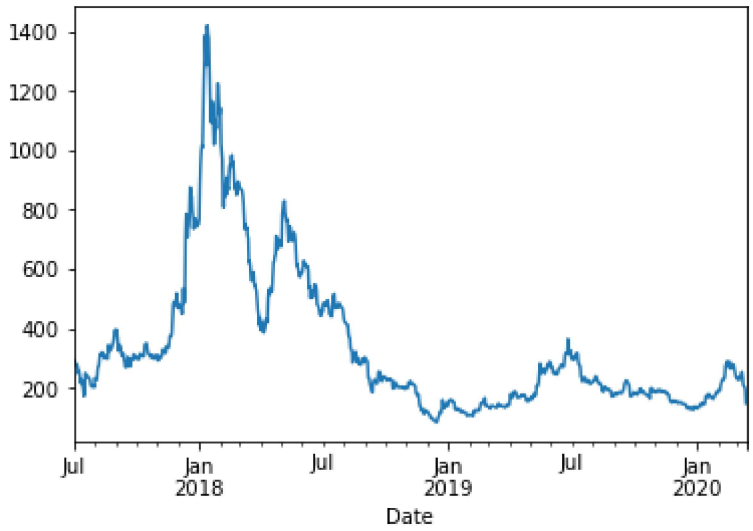

In [20]:
```
df.resample('W').mean()
```

Out[20]:

Date	Open	High	Low	Close	Volume
2017-07-02	268.066486	271.124595	264.819730	268.202162	2.185035e+06
2017-07-09	261.337024	262.872917	259.186190	261.062083	1.337349e+06
2017-07-16	196.193214	199.204405	192.722321	195.698393	2.986756e+06
2017-07-23	212.351429	215.779286	209.126310	212.783750	4.298593e+06
2017-07-30	203.496190	205.110357	201.714048	203.309524	1.581729e+06
...	...	...	...	...	...
2020-02-16	255.021667	257.255238	252.679762	255.198452	2.329087e+06
2020-02-23	265.220833	267.263690	262.948512	265.321905	1.826094e+06
2020-03-01	236.720536	238.697500	234.208750	236.373988	2.198762e+06
2020-03-08	229.923571	231.284583	228.373810	229.817619	1.628910e+06
2020-03-15	176.937521	179.979487	172.936239	176.332821	4.259828e+06

142 rows × 5 columns

In [22]:
```
df=df.resample('W').agg({'Close': 'mean', 'High': 'max', 'Low': 'min', 'Volume':
```

In [35]:
```
df
```

Out[35]:
```
<pandas.core.groupby.generic.DataFrameGroupBy object at 0x0000021D4EED28B0>
```

In [40]:
```
df.groupby('High').mean()
```

Out[40]:

High	Open	Low	Close	Volume
82.00	81.29	81.11	81.99	323206.25
82.13	81.92	81.80	81.89	528742.26
82.14	81.99	81.64	81.92	228693.57
82.32	82.32	81.00	81.29	395708.94
82.47	82.02	81.82	82.32	227169.03
...	...	...	...	...
1404.00	1395.00	1383.63	1392.24	9849624.19
1418.33	1410.00	1365.88	1382.56	23061241.98
1418.80	1382.56	1382.50	1418.61	15617541.24
1418.96	1388.99	1388.99	1410.00	27959588.92
1419.96	1418.61	1360.99	1395.00	11452917.63

16174 rows × 4 columns

In [30]:
```
df
```

Out[30]:

Date	Close	High	Low	Volume
2017-07-02	268.202162	293.73	253.23	8.084631e+07
2017-07-09	261.062083	285.00	231.25	2.246746e+08
2017-07-16	195.698393	240.33	130.26	5.017750e+08
2017-07-23	212.783750	249.40	153.25	7.221637e+08
2017-07-30	203.309524	229.99	178.03	2.657305e+08
...	...	...	...	...
2020-02-16	255.198452	290.00	216.31	3.912867e+08
2020-02-23	265.321905	287.13	242.36	3.067838e+08
2020-03-01	236.373988	278.13	209.26	3.693920e+08
2020-03-08	229.817619	253.01	196.00	2.736569e+08
2020-03-15	176.332821	208.65	90.00	4.983998e+08

142 rows × 4 columns

In [51]:
```
#Aggregate columns
df = pd.read_csv('D:/survey_results_public.csv', index_col='ResponseId')
schema_df = pd.read_csv('D:/survey_results_schema.csv')
```

In [54]:
```
df
```

Out[54]:

	Currency	CompTotal	CompFreq	LanguageHaveWorkedWith	
	NaN	NaN	NaN	NaN	
	CAD\tCanadian dollar	NaN	NaN	JavaScript;TypeScript	
	GBP\tPound sterling	32000.0	Yearly	C#;C++;HTML/CSS;JavaScript;Python	C#;C++
	ILS\tIsraeli new shekel	60000.0	Monthly	C#;JavaScript;SQL;TypeScript	
	USD\tUnited States dollar	NaN	NaN	C#;HTML/CSS;JavaScript;SQL;Swift;TypeScript	C#;Elixi
	...	...	...	...	
	USD\tUnited States dollar	60000.0	Yearly	Bash/Shell;Dart;JavaScript;PHP;Python;SQL;Type...	Bash/Shell;Go
	USD\tUnited States dollar	107000.0	Yearly	Bash/Shell;HTML/CSS;JavaScript;Python;SQL	
	USD\tUnited States dollar	NaN	NaN	HTML/CSS;JavaScript;PHP;Python;SQL	C#;HTML
	GBP\tPound sterling	58500.0	Yearly	C#;Delphi;VBA	

r	Currency	CompTotal	CompFreq	LanguageHaveWorkedWith
a	NaN	NaN	NaN	C#;JavaScript;Lua;PowerShell;SQL;TypeScript

In [46]:
```
pd.set_option('display.max_columns', 85)
pd.set_option('display.max_rows', 85)
```

In [49]:
```
df.head()
```

Out[49]:

LearnCodeCoursesCert	YearsCode	YearsCodePro	DevType	OrgSize	PurchaseInfluence	Buy
NaN	NaN	NaN	NaN	NaN	NaN	
NaN	NaN	NaN	NaN	NaN	NaN	
NaN	14	5	Data scientist or machine learning specialist;...	20 to 99 employees	I have some influence	
NaN	20	17	Developer, full-stack	100 to 499 employees	I have some influence	Othe
NaN	8	3	Developer, front-end;Developer, full-stack;Dev...	20 to 99 employees	I have some influence	St comn

In [56]:
```
df['ConvertedCompYearly'].median()
```

Out[56]: 67845.0

In [57]:
```
df.median()
```

Out[57]:
```
CompTotal                    77500.0
VCHostingPersonal use            NaN
VCHostingProfessional use        NaN
WorkExp                          8.0
ConvertedCompYearly          67845.0
dtype: float64
```

In [58]:
```
df.describe()
```

Out[58]:

	CompTotal	VCHostingPersonal use	VCHostingProfessional use	WorkExp	ConvertedCompYear
count	3.842200e+04	0.0	0.0	36769.000000	3.807100e+0
mean	2.342434e+52	NaN	NaN	10.242378	1.707613e+0
std	4.591478e+54	NaN	NaN	8.706850	7.814132e+0
min	0.000000e+00	NaN	NaN	0.000000	1.000000e+0
25%	3.000000e+04	NaN	NaN	4.000000	3.583200e+0
50%	7.750000e+04	NaN	NaN	8.000000	6.784500e+0
75%	1.540000e+05	NaN	NaN	15.000000	1.200000e+0
max	9.000000e+56	NaN	NaN	50.000000	5.000000e+0

In [59]:
```
df['ConvertedCompYearly'].count()
```

Out[59]:
```
38071
```

In [60]:
```
df['Country'].value_counts()
```

Out[60]:
```
United States of America                                   13543
India                                                       6639
Germany                                                     5395
United Kingdom of Great Britain and Northern Ireland        4190
Canada                                                      2490
                                                           ...
Seychelles                                                     1
Brunei Darussalam                                              1
Solomon Islands                                                1
Monaco                                                         1
Burkina Faso                                                   1
Name: Country, Length: 180, dtype: int64
```

In [61]:
```
country_grp = df.groupby(['Country'])
```

In [62]:
```
country_grp.get_group('India')
```

Out[62]:

ageHaveWorkedWith	LanguageWantToWorkWith	DatabaseHaveWorkedWith	
HP;Python;TypeScript	C;C#;C++;Elixir;Go;HTML/CSS;Java;JavaScript;Ko...	Cloud Firestore;MongoDB;Firebase Realtime Data...	
;Java;JavaScript;SQL	APL;Bash/Shell;Go;Python;TypeScript	MongoDB;MySQL	
C#	C++;JavaScript	MongoDB;MySQL	
cript;Kotlin;TypeScript	Groovy	Elasticsearch;PostgreSQL	
avaScript;Python;SQL	Bash/Shell;HTML/CSS;Java;JavaScript;Python;SQL	Oracle;PostgreSQL;Redis;SQLite	C
...	...	...	
Java;Python	Julia;Python	PostgreSQL;SQLite	
HTML/CSS;JavaScript	Go;HTML/CSS;Java;JavaScript;TypeScript	MongoDB;Firebase Realtime Database	
Script;Perl;PowerSh...	Bash/Shell;C#;HTML/CSS;JavaScript;Perl;PowerSh...	Microsoft SQL Server;MongoDB	
C;Python;SQL	C;C++	MySQL	
Go;Java;SQL	JavaScript;TypeScript	MongoDB;MySQL	

In [64]:
```
filt = df['Country'] == 'India'
df.loc[filt]['OfficeStackSyncHaveWorkedWith'].value_counts()
```

Out[64]:
```
Microsoft Teams                                                    680
Microsoft Teams;Zoom                                               504
Slack;Zoom                                                         468
Zoom                                                               419
Slack                                                              373
                                                                  ...
Wickr;Zoom                                                           1
Unify Circuit;Zoom                                                   1
Cisco Webex Teams;Mattermost;Microsoft Teams;Rocketchat;Slack;Zoom   1
Microsoft Teams;RingCentral;Slack;Zoom                               1
Microsoft Teams;RingCentral;Symphony;Zoom                            1
Name: OfficeStackSyncHaveWorkedWith, Length: 144, dtype: int64
```

In [66]:
```
filt = df['Country'] == 'India'
df.loc[filt]['LanguageWantToWorkWith'].str.contains('Python').sum()
```

Out[66]:
```
3094
```

In [70]:
```
country_uses_python=country_grp['LanguageWantToWorkWith'].apply(lambda x: x.str.
```

In [68]:
```
country_respondents = df['Country'].value_counts()
country_respondents
```

Out[68]:
```
United States of America                              13543
India                                                  6639
Germany                                                5395
United Kingdom of Great Britain and Northern Ireland   4190
Canada                                                 2490
                                                      ...
Seychelles                                                1
Brunei Darussalam                                         1
Solomon Islands                                           1
Monaco                                                    1
Burkina Faso                                              1
Name: Country, Length: 180, dtype: int64
```

In [71]:
```
python_df = pd.concat([country_respondents, country_uses_python], axis='columns',
python_df
```

Out[71]:

	Country	LanguageWantToWorkWith
United States of America	13543	5656
India	6639	3094
Germany	5395	2212
United Kingdom of Great Britain and Northern Ireland	4190	1594
Canada	2490	1000
...	...	...
Seychelles	1	0
Brunei Darussalam	1	1
Solomon Islands	1	1
Monaco	1	1
Burkina Faso	1	1

180 rows × 2 columns

In [72]:
```
python_df.rename(columns={'Country': 'NumRespondents', 'LanguageWantToWorkWith':
```

In [73]:
```
python_df
```

Out[73]:

	NumRespondents	NumKnowsPython
United States of America	13543	5656
India	6639	3094
Germany	5395	2212
United Kingdom of Great Britain and Northern Ireland	4190	1594
Canada	2490	1000
...	...	...
Seychelles	1	0
Brunei Darussalam	1	1
Solomon Islands	1	1
Monaco	1	1
Burkina Faso	1	1

180 rows × 2 columns

In [74]:

```
NumRespondents  NumKnowsPython
United States   20949   10083
India   9061    3105
Germany 5866    2451
United Kingdom  5737    2384
Canada  3395    1558
... ... ...
Dominica    1   1
Tonga   1   0
Sao Tome and Principe   1   1
Saint Kitts and Nevis   1   0
Brunei Darussalam   1   0
179 rows × 2 columns

python_df['PctKnowsPython'] = (python_df['NumKnowsPython']/python_df['NumRespondе
python_df
```

```
  File "<ipython-input-74-b9056b2d7647>", line 1
    NumRespondents      NumKnowsPython
                        ^
SyntaxError: invalid syntax
```

In [75]:

```
python_df['PctKnowsPython'] = (python_df['NumKnowsPython']/python_df['NumRespondе
python_df
```

Out[75]:

	NumRespondents	NumKnowsPython	PctKnowsPython
United States of America	13543	5656	41.763273
India	6639	3094	46.603404
Germany	5395	2212	41.000927
United Kingdom of Great Britain and Northern Ireland	4190	1594	38.042959
Canada	2490	1000	40.160643
...	...	...	...
Seychelles	1	0	0.000000
Brunei Darussalam	1	1	100.000000
Solomon Islands	1	1	100.000000
Monaco	1	1	100.000000
Burkina Faso	1	1	100.000000

180 rows × 3 columns

In [76]:
```
python_df.loc['Japan']
```

Out[76]:
```
NumRespondents    333.000000
NumKnowsPython    128.000000
PctKnowsPython     38.438438
Name: Japan, dtype: float64
```

In [78]:
```
#Sorting values
df.sort_values(by=['Country', 'ConvertedCompYearly'], ascending=[True, False], ir
```

In [80]:
```
df.head()
```

Out[80]:

ledge_3	Knowledge_4	Knowledge_5	Knowledge_6	Knowledge_7	Frequency_1	Frequency_2	Frequ
NaN	NaN	NaN	NaN	NaN	NaN	NaN	
NaN	NaN	NaN	NaN	NaN	NaN	NaN	
)isagree	Agree	Agree	Strongly agree	Neither agree nor disagree	1-2 times a week	1-2 times a week	1-2
er agree disagree	Neither agree nor disagree	Neither agree nor disagree	Neither agree nor disagree	Neither agree nor disagree	1-2 times a week	1-2 times a week	1-2
Strongly agree	Strongly agree	Strongly agree	Agree	Agree	Never	1-2 times a week	1-2

In [81]:
```
df['ConvertedCompYearly'].nlargest(10)
```

Out[81]:
```
ResponseId
40305    50000000.0
202      44790396.0
62027    35000000.0
70523    32500000.0
61044    28853768.0
18923    22500000.0
62224    22500000.0
66496    18000000.0
1291     15000000.0
24164    15000000.0
Name: ConvertedCompYearly, dtype: float64
```

In [82]:
```
df.nsmallest(10, 'ConvertedCompYearly')
```

Out[82]:

3	TimeSearching	TimeAnswering	Onboarding	ProfessionalTech	TrueFalse_1	TrueFalse_
a <	15-30 minutes a day	30-60 minutes a day	Very long	Innersource initiative;DevOps function;Microse...	Yes	Y
a <	60-120 minutes a day	Over 120 minutes a day	Just right	DevOps function;Microservices;Developer portal...	Yes	Y
a <	15-30 minutes a day	30-60 minutes a day	Somewhat short	DevOps function;Microservices;Developer portal...	Yes	Y
J	NaN	NaN	NaN	NaN	NaN	Na
r	30-60 minutes a day	15-30 minutes a day	Very short	None of these	Yes	Y
a <	15-30 minutes a day	15-30 minutes a day	Just right	Automated testing	Yes	N
J	NaN	NaN	NaN	NaN	NaN	Na
r	15-30 minutes a day	15-30 minutes a day	Just right	Innersource initiative;DevOps function;Microse...	Yes	Y
J	NaN	NaN	NaN	NaN	NaN	Na
r	30-60 minutes a day	Less than 15 minutes a day	Just right	Continuous integration (CI) and (more often) c...	No	N

In [83]:
```
#Add remove rows and columns
people = {
    'first': ['Corey', 'Jane', 'John'],
    'last': ['Schafer', 'Doe', 'Doe'],
    'email': ['CoreyMSchafer@gmail.com', 'JaneDoe@email.com', 'JohnDoe@email.com'
}
```

In [84]:
```
df = pd.DataFrame(people)
```

In [85]:
```
df
```

Out[85]:

	first	last	email
0	Corey	Schafer	CoreyMSchafer@gmail.com
1	Jane	Doe	JaneDoe@email.com
2	John	Doe	JohnDoe@email.com

In [86]:
```
df['first'] + ' ' + df['last']
```

Out[86]:
```
0    Corey Schafer
1         Jane Doe
2         John Doe
dtype: object
```

In [87]:
```
df['full_name'] = df['first'] + ' ' + df['last']
```

In [88]:
```
df
```

Out[88]:

	first	last	email	full_name
0	Corey	Schafer	CoreyMSchafer@gmail.com	Corey Schafer
1	Jane	Doe	JaneDoe@email.com	Jane Doe
2	John	Doe	JohnDoe@email.com	John Doe

In [89]:
```
df.drop(columns=['first', 'last'], inplace=True)
```

In [90]:
```
df
```

Out[90]:

	email	full_name
0	CoreyMSchafer@gmail.com	Corey Schafer
1	JaneDoe@email.com	Jane Doe
2	JohnDoe@email.com	John Doe

In [91]:
```
df['full_name'].str.split(' ', expand=True)
```

Out[91]:

	0	1
0	Corey	Schafer
1	Jane	Doe
2	John	Doe

In [92]:
```
df[['first', 'last']] = df['full_name'].str.split(' ', expand=True)
```

In [93]:
```
df
```

Out[93]:

	email	full_name	first	last
0	CoreyMSchafer@gmail.com	Corey Schafer	Corey	Schafer
1	JaneDoe@email.com	Jane Doe	Jane	Doe
2	JohnDoe@email.com	John Doe	John	Doe

In [94]:
```
df.append({'first': 'Tony'}, ignore_index=True)
```

Out[94]:

	email	full_name	first	last
0	CoreyMSchafer@gmail.com	Corey Schafer	Corey	Schafer
1	JaneDoe@email.com	Jane Doe	Jane	Doe
2	JohnDoe@email.com	John Doe	John	Doe
3	NaN	NaN	Tony	NaN

In [98]:
```
drop=df.drop(index=2)
```

In [99]:
```
drop
```

Out[99]:

	email	full_name	first	last
0	CoreyMSchafer@gmail.com	Corey Schafer	Corey	Schafer
1	JaneDoe@email.com	Jane Doe	Jane	Doe

In [100]:
```
df
```

Out[100]:

	email	full_name	first	last
0	CoreyMSchafer@gmail.com	Corey Schafer	Corey	Schafer
1	JaneDoe@email.com	Jane Doe	Jane	Doe
2	JohnDoe@email.com	John Doe	John	Doe

In [103]:
```
filt=df['last']=='Doe'
df.drop(index=df[filt].index)
```

Out[103]:

	email	full_name	first	last
0	CoreyMSchafer@gmail.com	Corey Schafer	Corey	Schafer

In [106]:
```
people = {
    'first': ['Tony', 'Steve'],
    'last': ['Stark', 'Rogers'],
    'email': ['IronMan@avenge.com', 'Cap@avenge.com']
}
df2 = pd.DataFrame(people)
df2
```

Out[106]:

	first	last	email
0	Tony	Stark	IronMan@avenge.com
1	Steve	Rogers	Cap@avenge.com

In [105]:
```
df.append(df2, ignore_index=True, sort=False)
```

Out[105]:

	email	full_name	first	last
0	CoreyMSchafer@gmail.com	Corey Schafer	Corey	Schafer
1	JaneDoe@email.com	Jane Doe	Jane	Doe
2	JohnDoe@email.com	John Doe	John	Doe
3	IronMan@avenge.com	NaN	Tony	Stark
4	Cap@avenge.com	NaN	Steve	Rogers

In [109]:
```
#Renaming columns
#Aggregate columns
df = pd.read_csv('D:/survey_results_public.csv', index_col='ResponseId')
schema_df = pd.read_csv('D:/survey_results_schema.csv')
```

In [110]:
```
df
```

Out[110]:

ResponseId	MainBranch	Employment	RemoteWork	CodingActivities	EdLevel	LearnCode	
1	None of these	NaN	NaN	NaN	NaN	NaN	
2	I am a developer by profession	Employed, full-time	Fully remote	Hobby;Contribute to open-source projects	NaN	NaN	
3	I am not primarily a developer, but I write co...	Employed, full-time	Hybrid (some remote, some in-person)	Hobby	Master's degree (M.A., M.S., M.Eng., MBA, etc.)	Books / Physical media;Friend or family member...	do
4	I am a developer by profession	Employed, full-time	Fully remote	I don't code outside of work	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Books / Physical media;School (i.e., Universit...	
5	I am a developer by profession	Employed, full-time	Hybrid (some remote, some in-person)	Hobby	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Other online resources (e.g., videos, blogs, f...	
...	...	...	...	...	...	...	
73264	I am a developer by profession	Employed, full-time	Fully remote	Freelance/contract work	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Books / Physical media;Other online resources ...	
73265	I am a developer by profession	Employed, full-time	Full in-person	Hobby	Master's degree (M.A., M.S., M.Eng., MBA, etc.)	Other online resources (e.g., videos, blogs, f...	
73266	I am not primarily a developer, but I write co...	Employed, full-time	Hybrid (some remote, some in-person)	Hobby;School or academic work	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Books / Physical media;Other online resources ...	
73267	I am a developer by profession	Employed, full-time	Hybrid (some remote, some in-person)	Hobby	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Books / Physical media;On the job training	

ResponseId	MainBranch	Employment	RemoteWork	CodingActivities	EdLevel	LearnCode	
73268	I used to be a developer by profession, but no...	Independent contractor, freelancer, or self-em...	Fully remote	Hobby;Contribute to open-source projects;Boots...	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Books / Physical media;Friend or family member...	do

73268 rows × 78 columns

In [111]:
```
pd.set_option('display.max_columns', 85)
pd.set_option('display.max_rows', 85)
```

In [112]:
```
df.head()
```

Out[112]:

ResponseId	MainBranch	Employment	RemoteWork	CodingActivities	EdLevel	LearnCode	
1	None of these	NaN	NaN	NaN	NaN	NaN	
2	I am a developer by profession	Employed, full-time	Fully remote	Hobby;Contribute to open-source projects	NaN	NaN	
3	I am not primarily a developer, but I write co...	Employed, full-time	Hybrid (some remote, some in-person)	Hobby	Master's degree (M.A., M.S., M.Eng., MBA, etc.)	Books / Physical media;Friend or family member...	doc
4	I am a developer by profession	Employed, full-time	Fully remote	I don't code outside of work	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Books / Physical media;School (i.e., Universit...	
5	I am a developer by profession	Employed, full-time	Hybrid (some remote, some in-person)	Hobby	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Other online resources (e.g., videos, blogs, f...	

In [113]:
```
df.rename(columns={'ConvertedCompYearly': 'SalaryUSD'}, inplace=True)
```

In [116]:
```
df['SalaryUSD']
```

Out[116]:
```
ResponseId
1             NaN
2             NaN
3         40205.0
4        215232.0
5             NaN
             ...
73264         NaN
73265         NaN
73266         NaN
73267         NaN
73268         NaN
Name: SalaryUSD, Length: 73268, dtype: float64
```

In [ ]:
```

```